

Supercompilation by Evaluation

Max Bolingbroke (University of Cambridge)
Simon Peyton Jones (Microsoft Research)

To supercompile, evaluate!

```
let fact = λn. if n == 0 then 1 else n * fact (n - 1)
in fact 1
```



1

- **Supercompilation complete!**
- That was almost too easy...
- It is not always so straightforward - it gets more interesting, as we will see

Problem one: free variables

```
let inc = λx. x+1
    map = λf xs. case xs of [] -> []
                                   (y:ys) -> f y : map f ys
in map inc zs
```



```
let inc = ...; map = ...
in case zs of [] -> []
              (y:ys) -> inc y : map inc ys
```

- Supercompilers need to deal with terms with **missing information**
- We **can't know** at compile time what list we are mapping over. What to do?

Solution one: split around the free variable




```
let inc = ...; map = ...  
in case zs of []      -> []  
             (y:ys) -> inc y : map inc ys
```

- To supercompile when **stuck**, “split” to supercompile some subcomponents:

“**Residual**” shell

**Recursively
supercompiled**

expressions: result is
plugged into the holes ●



Problem two: values containing fields

```
let inc = ...; map = ...  
in map inc [1, 2, 3]
```



```
let inc = ...; map = ...  
in inc 1 : map inc [2, 3]
```

- Another kind of problem can arise when **call-by-{name,need}** reduction reaches a value
- Due to **laziness**, evaluation by itself can't touch e.g. the `inc 1` subterm... but we'd really like for the supercompiler to reduce it to 2


Solution two: split around the structure of the value

```
let inc = ...; map = ...  
in inc 1 : map inc [2, 3]
```

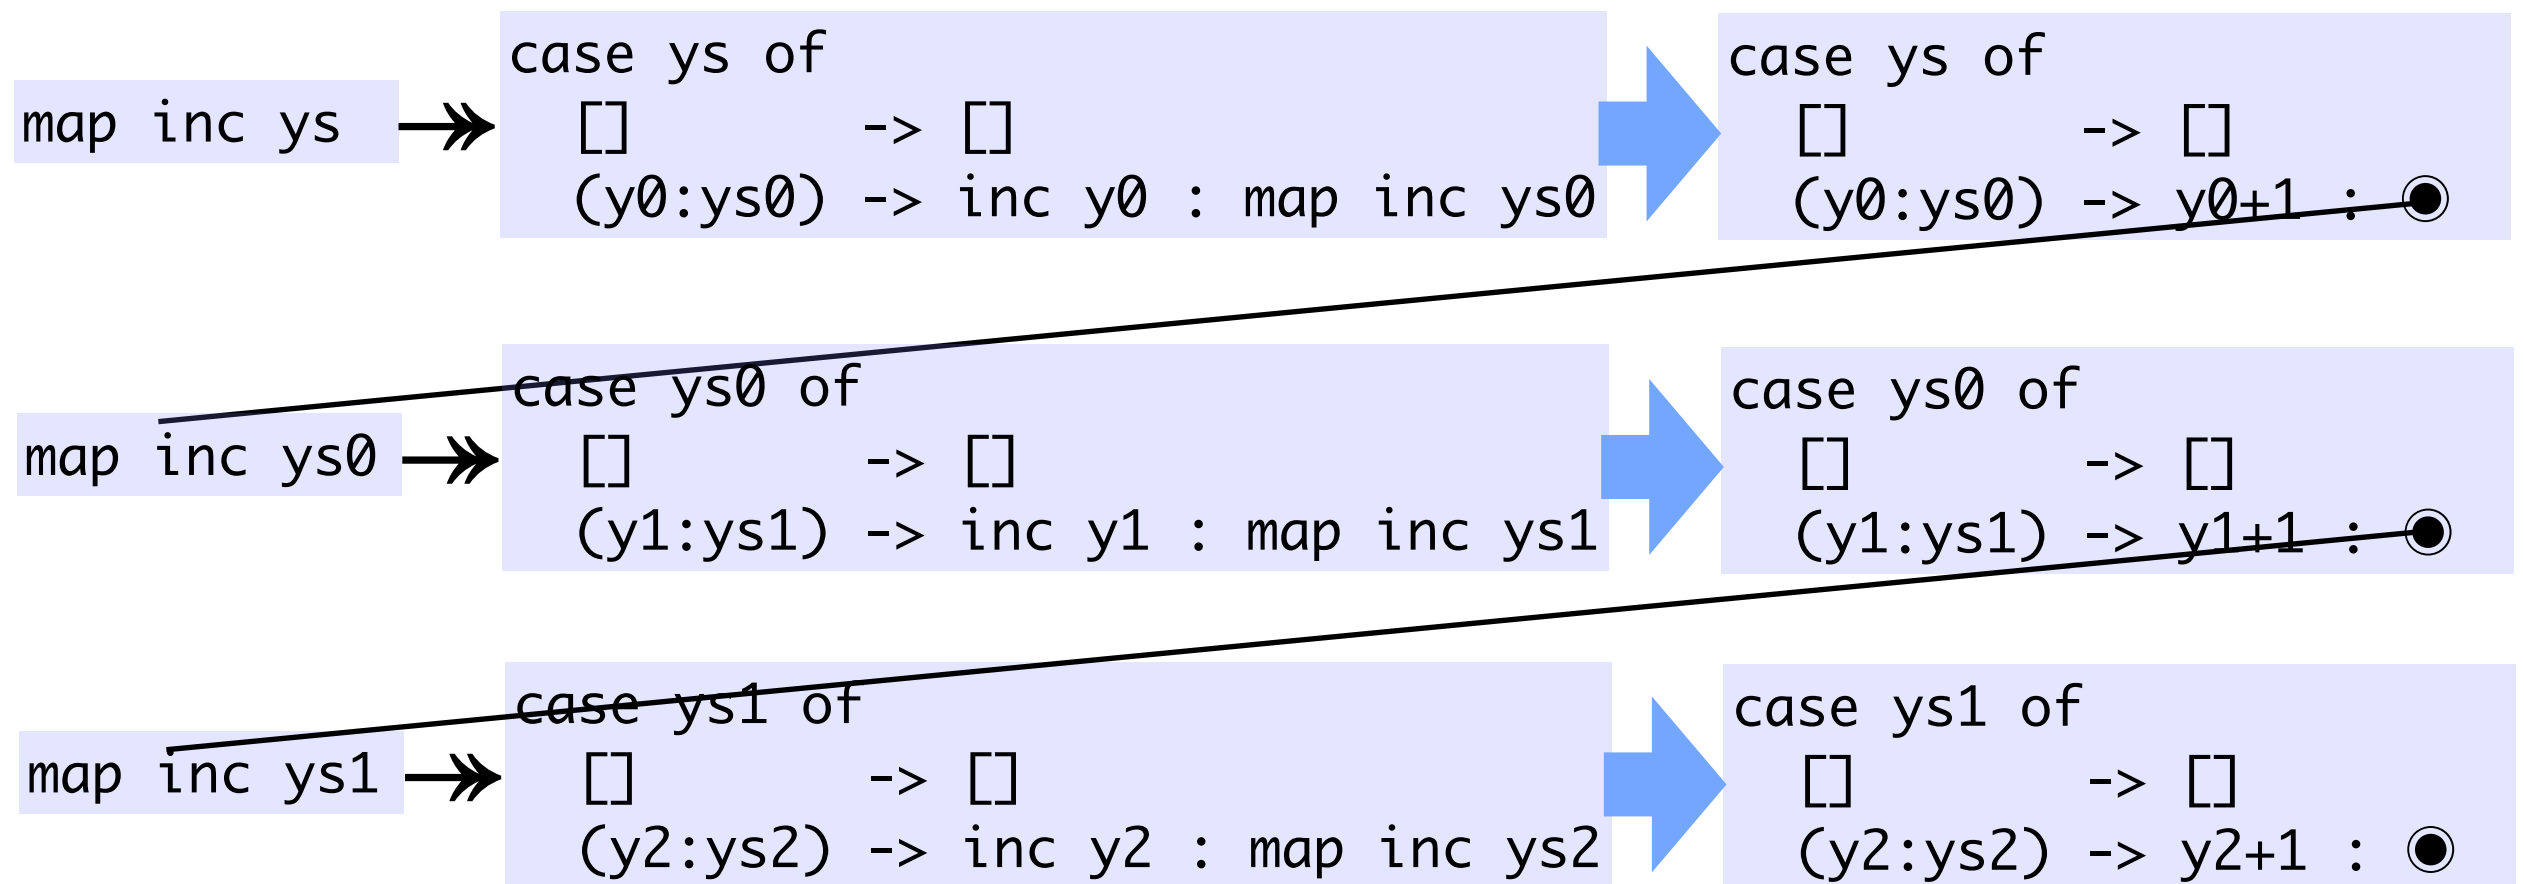
- Play exactly the same game as before: “split” the term to find some subcomponents amenable to driving

Residual shell

**Recursively
supercompiled**
expressions: result is
plugged into the holes ●



Problem three: infinite programs



- ... and so on
- It looks like the supercompiler will produce an **infinite output program**, and thus **never terminate!**

Infinite programs?

map inc ys

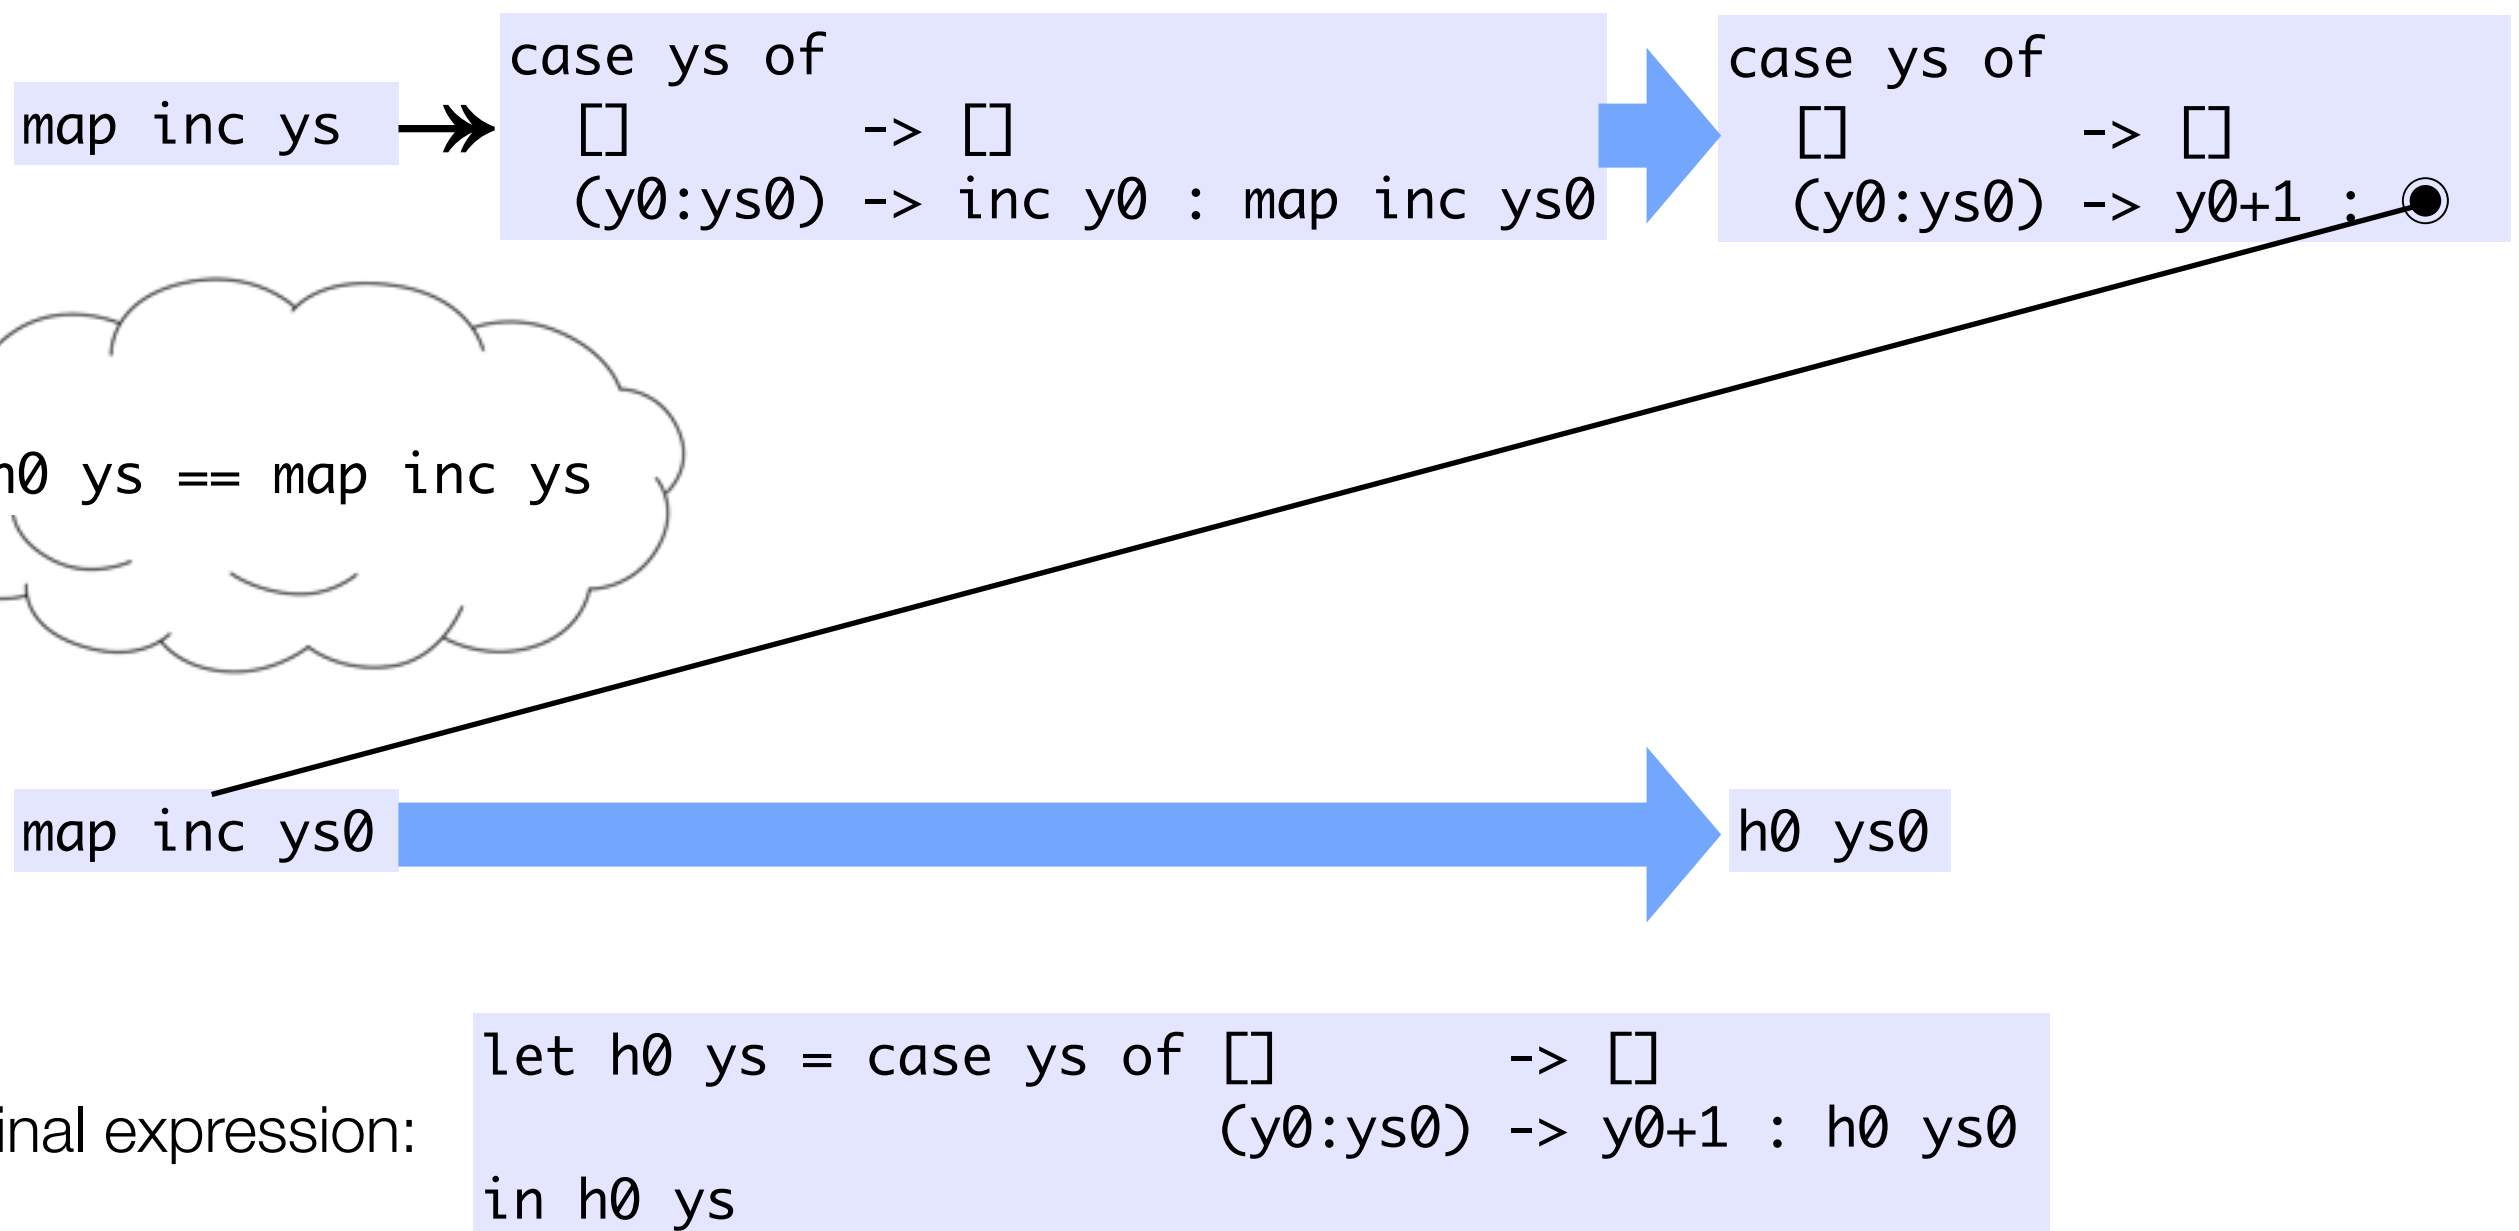


```
case ys of
  []      -> []
  (y0:ys0) -> y0+1 : case ys0 of
    []      -> []
    (y1:ys1) -> y1+1 : case ys1 of
      []      -> []
      (y2:ys2) -> case ys2 of ...
```

- Maybe not what we had in mind...

Solution three: memoisation

- If we **memoise** the supercompiler by **let-binding** supercompiled terms, we can get **finite output programs**:



Supercompilation \approx Memoisation
+ Evaluation

... but no other supercompiler does
evaluation right

Existing work

- Existing supercompilers are based on evaluators only in an **ad-hoc** way
 - Typically based on a **driving function** which intermingles the evaluation and splitting steps
- Supero 2010 is slightly better: it separates out a **simplifier** :: Term -> Maybe Term
 - Still not quite an evaluator - doesn't follow standard evaluation order
 - Does not terminate on some terms...
- The ad-hoc evaluators used so far are **incapable** of reducing some call-by-need terms (they don't like local **recursive** let bindings)

Recursive let

```
let ones = 1 : ones  
in head ones
```

```
let repeat =  $\lambda x \rightarrow$  let xs = x : xs  
                    in xs  
in repeat 1
```

```
map ( $\backslash x \rightarrow x + 1$ ) (repeat 1)
```



```
let twos = 2 : twos in twos
```

- My feeling is that you can't claim to have a supercompiler for a call by need language if you leave some features unimplemented like this

The solution

- Clean semantics for evaluating languages that include recursive `let` are well known [Sestoft 1993, Launchbury 1993, Wadler 1998]
- Stop trying to reinvent the wheel - just take one of those **unmodified** and build our implementation on top

`step :: State -> Maybe State`

The state of an abstract machine for evaluating the language

Sestoft's machine

type State = (Heap, Term, Stack)

An environment mapping
variables to Terms

The focus of evaluation:
we want to reduce this to a
value

What is around the focus -
what to do with the value
once we have it

```
< ε
| let t = True
  in case (\x -> x) t of
    True  -> 1
    False -> 2
| ε >
```

```
< t = True
| case (\x -> x) t of
  True  -> 1
  False -> 2
| ε >
```

```
< t = True
| (\x -> x) t
| case • of
  True  -> 1
  False -> 2 >
```

```
< t = True
| True
| case • of
  True  -> 1
  False -> 2 >
```

```
< t = True
| t
| case • of
  True  -> 1
  False -> 2 >
```

```
< t = True
| \x -> x
| • t,
  case • of
    True  -> 1
    False -> 2 >
```

Sestoft's machine vs Mitchell

```
let x1 = e1
    ...
    xn = en
in xi
```

```
([x1 |-> e1, ..., xn |-> en], xi, [])
```

Sestoft's machine vs Jonsson

$$\mathcal{R} ::= [] \mid \mathcal{R} \bar{e} \mid \text{case } \mathcal{R} \text{ of } \{p_i \rightarrow e_i\} \mid \mathcal{R} \oplus e \mid e \oplus \mathcal{R} \mid \mathcal{R}.l$$

κ	$::=$	update x	Update frame
		$\bullet x$	Apply to function value
		case \bullet of $\overline{\alpha} \rightarrow \overline{e}$	Scrutinise value
		$\bullet \otimes e$	Apply first value to primop
		$v \otimes \bullet$	Apply second value to primop

Apotheosis

~~scp :: Term -> ScpM Term~~

scp :: State -> ScpM Term

See paper for the gory details



Benefits of using a real evaluator

- **Modularity:** the issues of evaluation are separated from those specific to supercompilation (memoisation, splitting...)
- **Generality:** it is easy to swap in *alternative* evaluation strategies (call-by-value, garbage collecting, let-speculating...)
- **Optimisation:** we can supercompile programs that no existing call-by-need supercompiler can deal with:

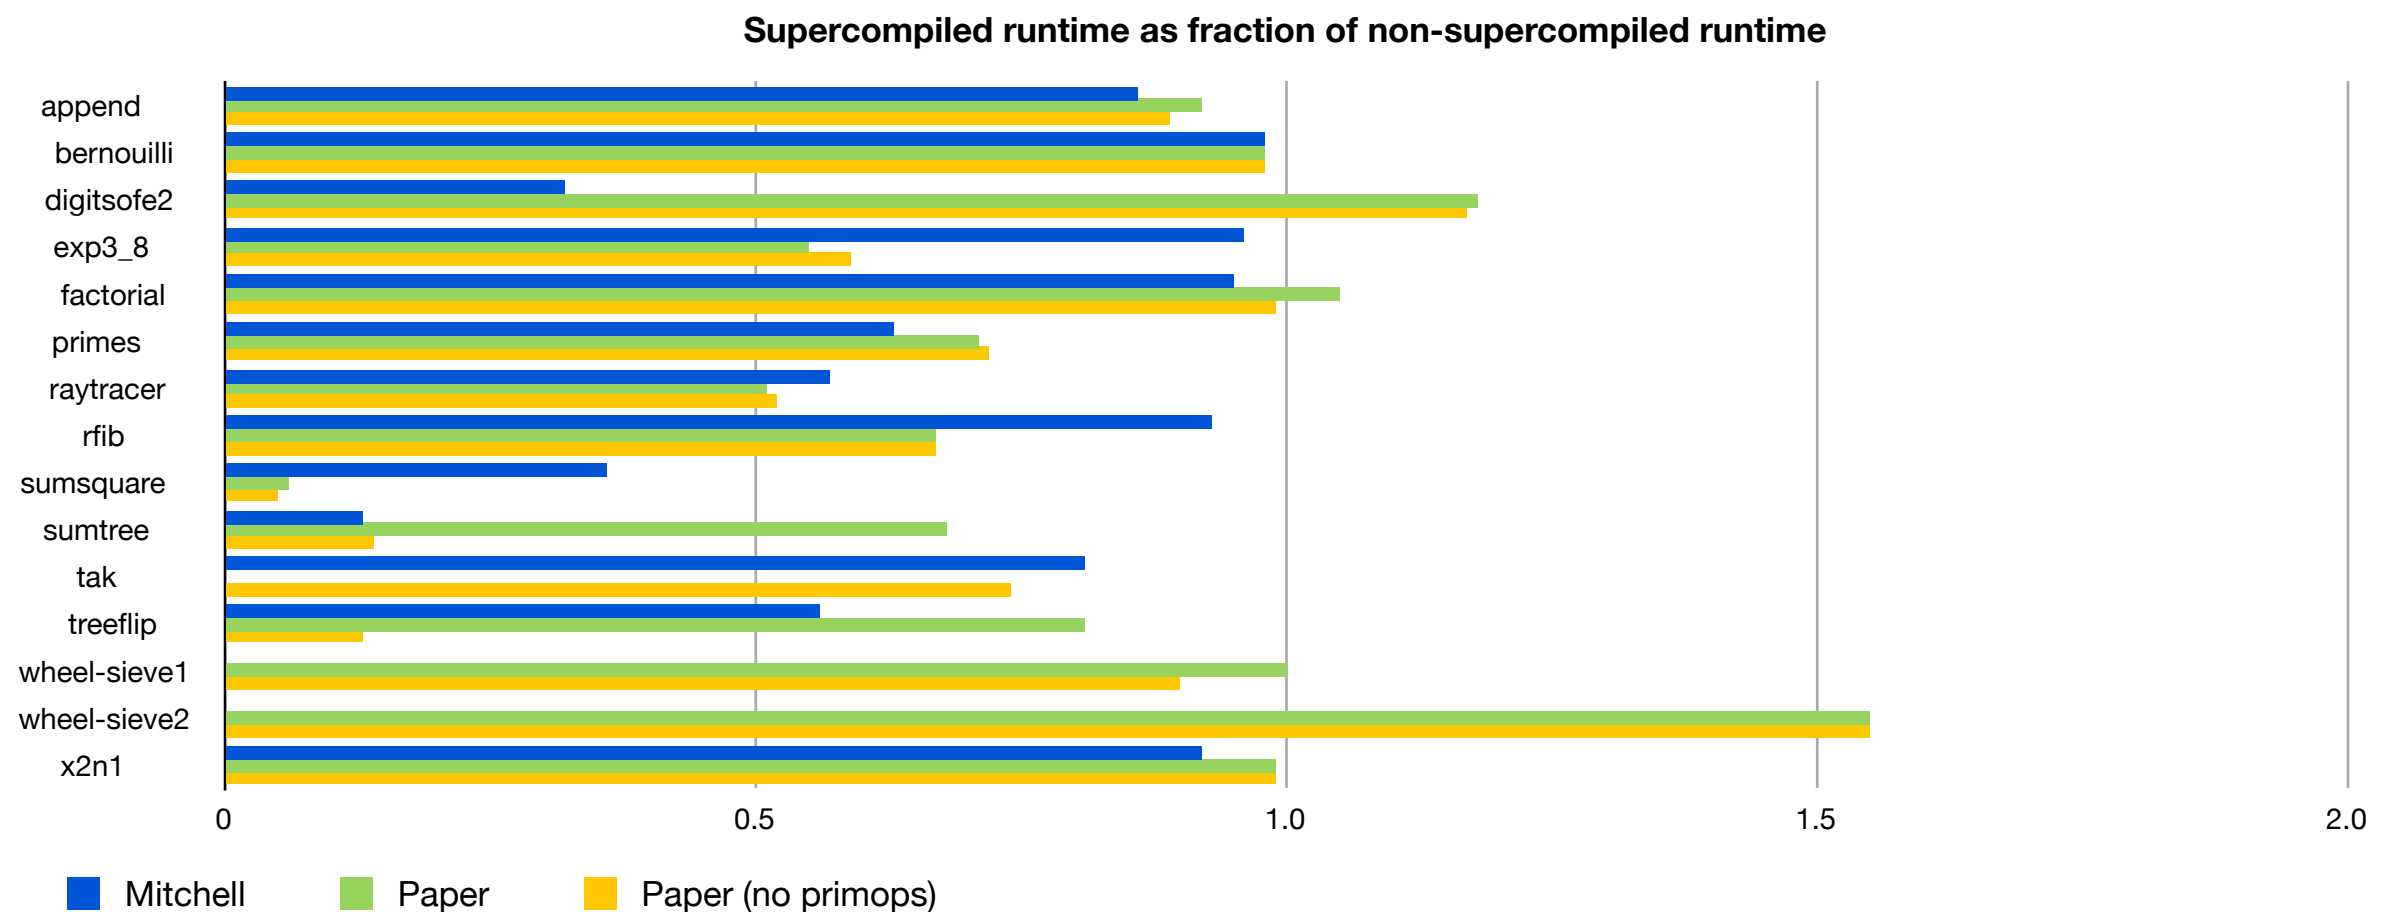
```
map (\x -> x + 1) (repeat 1)
```



```
let twos = 2 : twos in twos
```

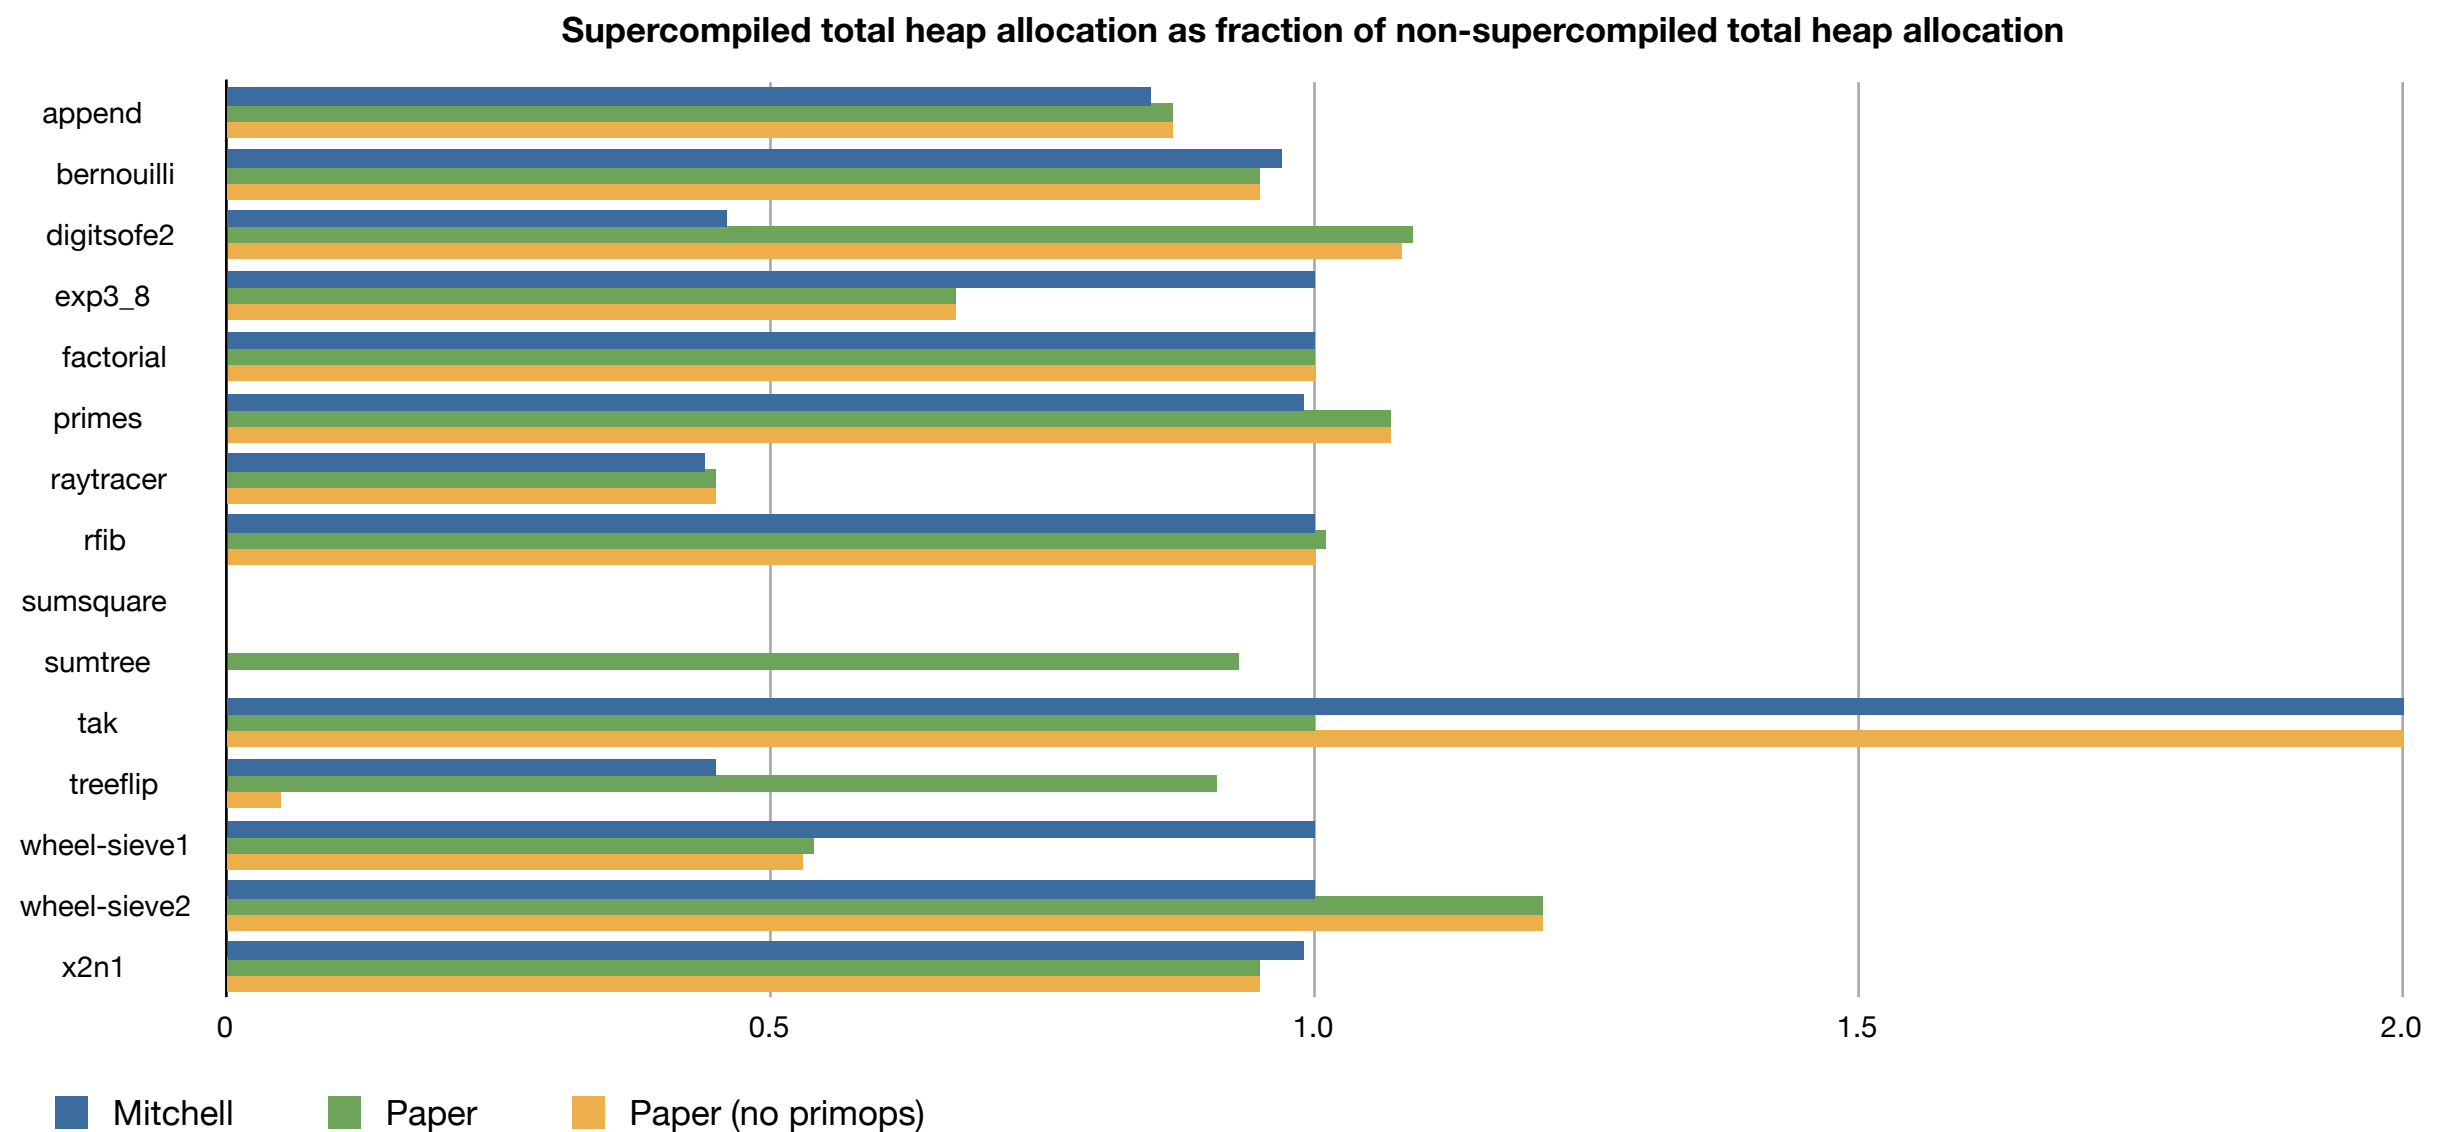
Results

- Our implementation, CHSC, is a preprocessor which produces supercompiled Haskell code for later standard compilation by GHC



- On benchmarks we can both supercompile, average runtime reduction is ~30% for both Mitchell and our supercompiler (without primops)


Results



- Heap allocation totally eliminated for e.g. sumsquare
- 600x worse (Mitchell) or 19,000x worse (us) for e.g. tak

Results

Supercompiled term size as fraction of non-supercompiled term size



- On programs we both supercompile, the average size increase is ~2.5x (Mitchell) / ~15x (us)

Related work

	Language	Recursive let?
SCP4 [Nemytykh]	Refal (CBV)	✗
Timber supercompiler [Jonsson and Nordlander]	Lambda calculus (CBV)	✗
Supero 2010 [Mitchell]	Lambda calculus (Lazy)	✗
HOSC [Klyuchnikov and Romanenko]	Lambda calculus (CBN)	✓
CHSC [Bolingbroke and Peyton Jones]	Lambda calculus (Lazy)	✓

Future work

- **Better optimisation** for output programs
 - In particular, use **generalisation** to prevent specialisation on accumulators
- Make supercompilation **faster**
 - Several imaginary suite benchmarks **cannot be supercompiled in reasonable amounts of time** by any existing system
 - Even some of the examples that did supercompile successfully are annoyingly slow: wheel-sieve1 takes 22s to supercompile

Conclusion

- Takeaway: supercompilers are better when built on top of **real evaluators**
 - **Simpler** to implement and read
 - More **powerful**

Thanks!

<http://www.github.com/batterseapower/chsc>