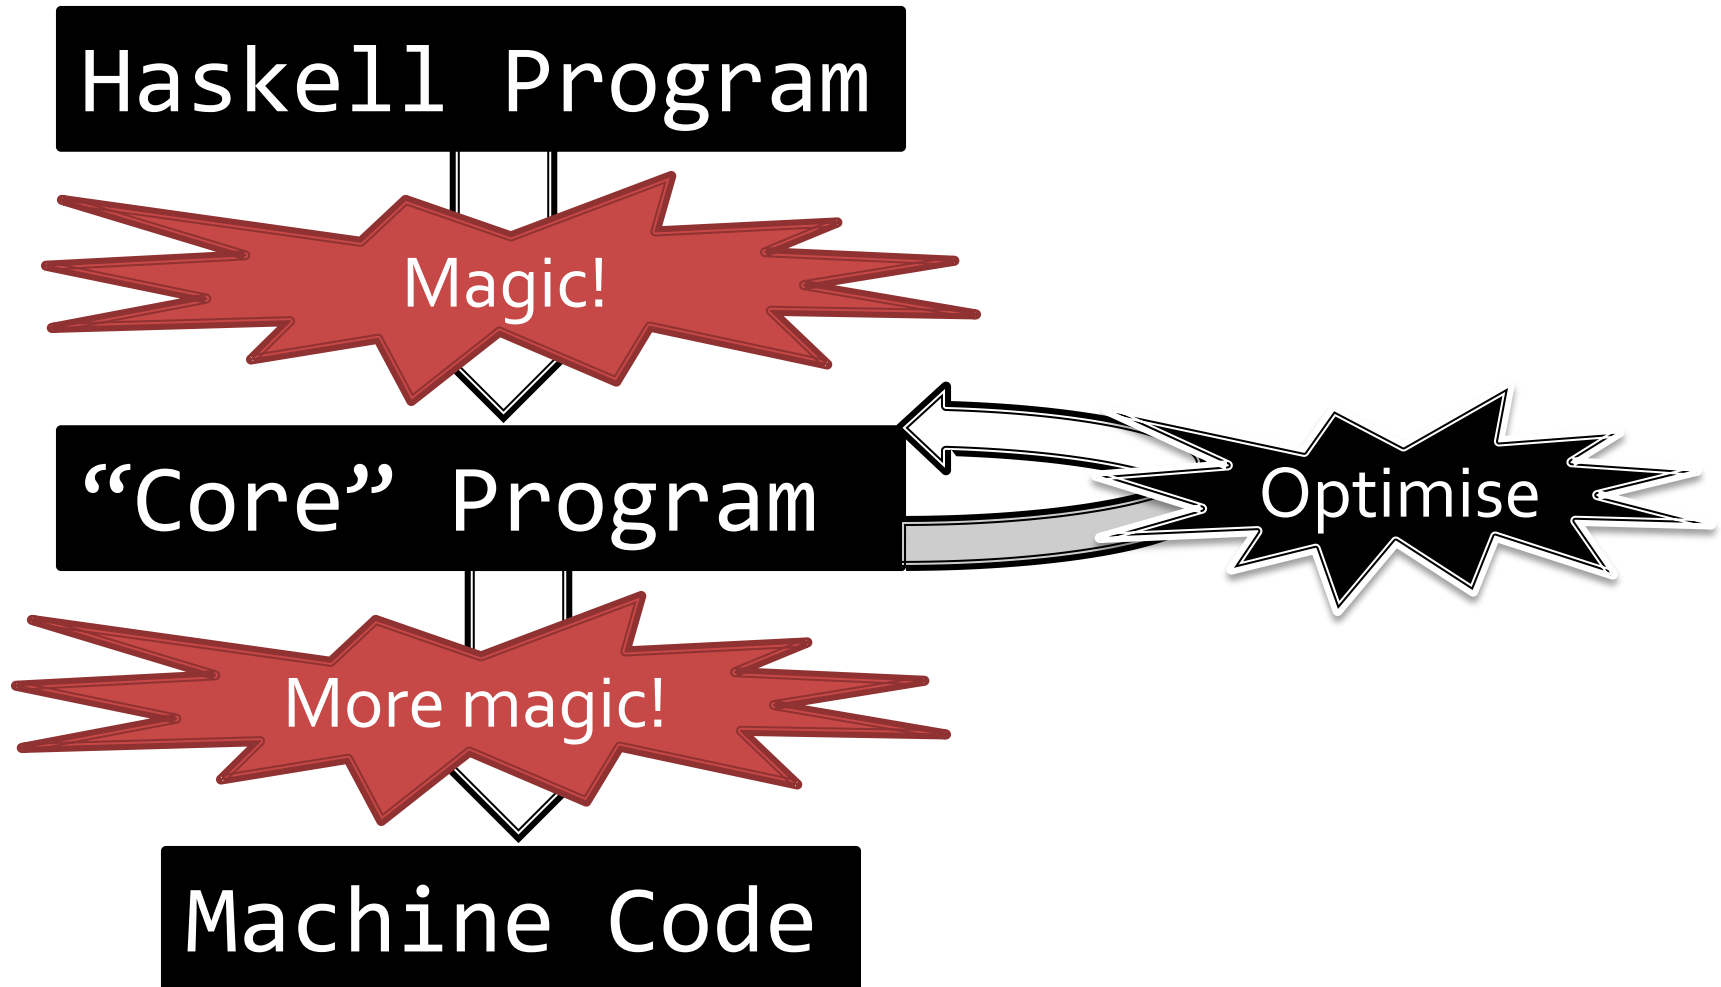


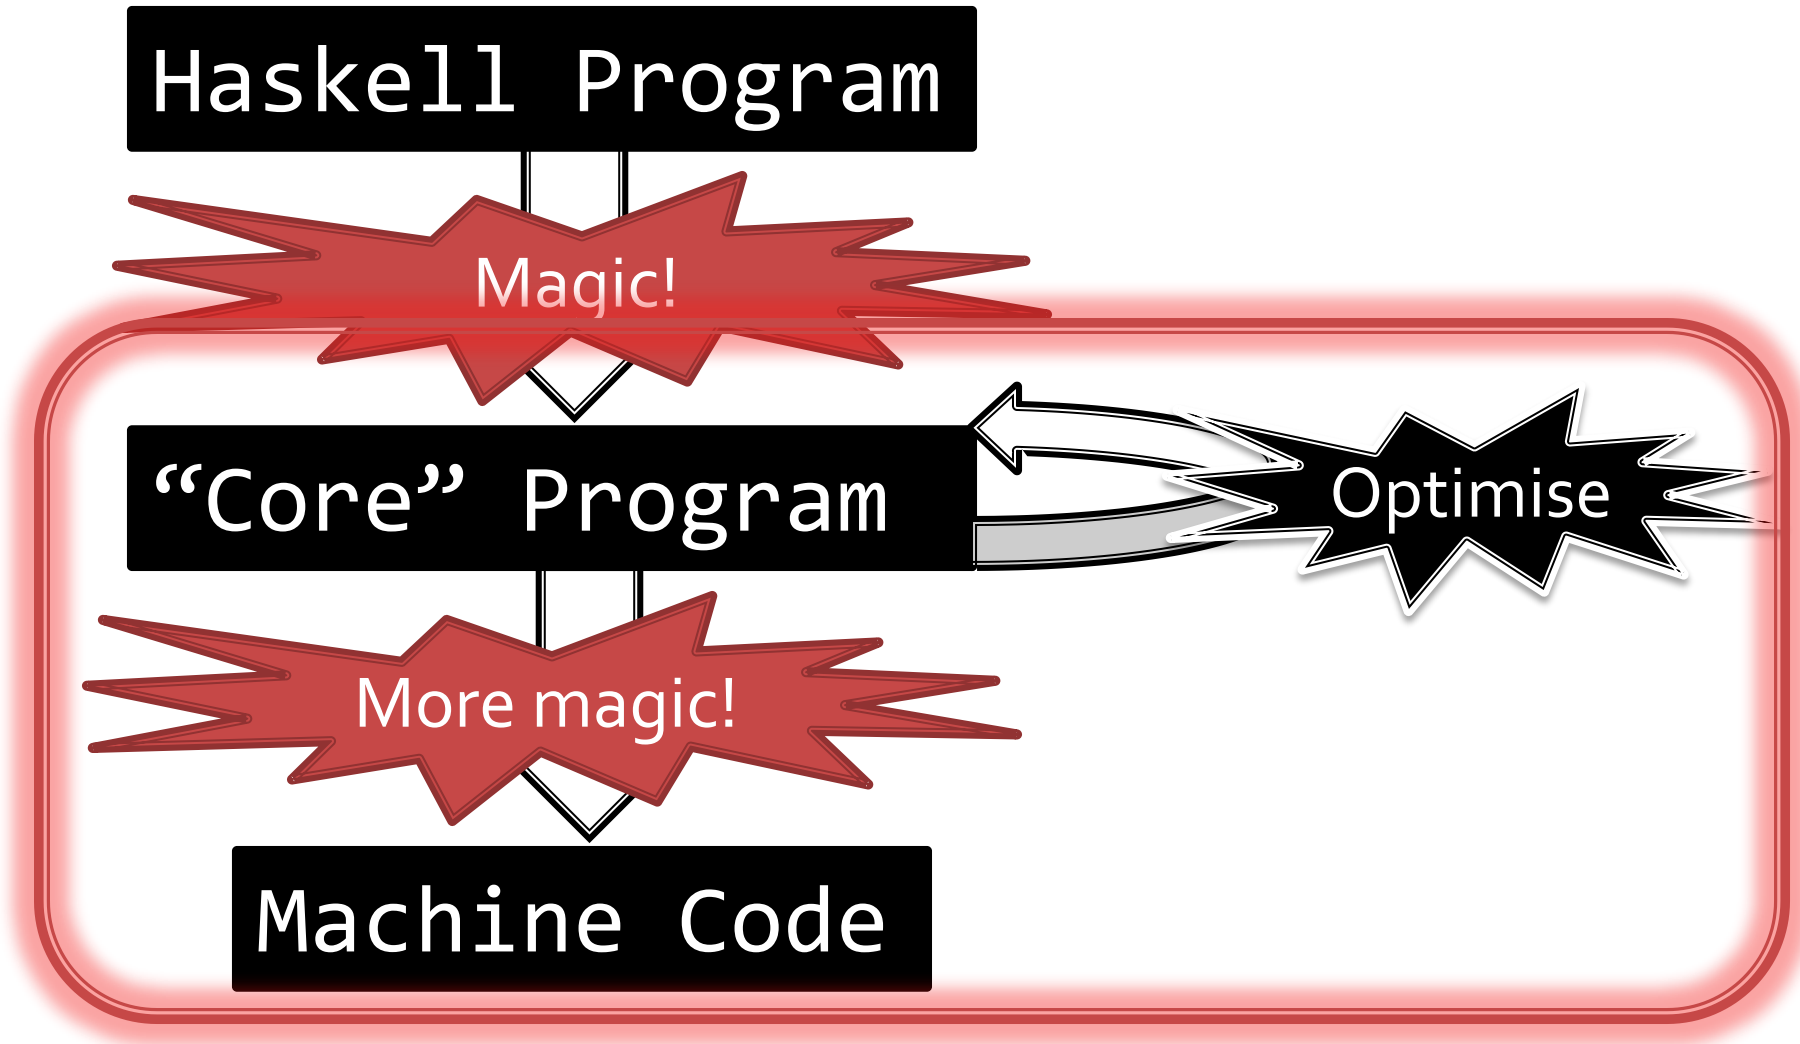
Max Bolingbroke
Simon Peyton Jones

Types Are Calling Conventions

The Glasgow Haskell Compiler



The Glasgow Haskell Compiler



How Many Parameters? 1

How many parameters can you supply to a function of this type?

```
f :: Int -> Int -> Int
```

As far as the user is concerned, the answer is obviously two:

```
f 1 2
```

How Many Parameters? 2

However, the compiler actually makes a finer grained distinction:

How Many Parameters? 2

However, the compiler actually makes a finer grained distinction:

Two: `f2 = \y -> \z -> let x = fib 10 in x + y + z`

How Many Parameters? 2

However, the compiler actually makes a finer grained distinction:

Two: `f2 = \y -> \z -> let x = fib 10 in x + y + z`

One: `f1 = \y -> let x = fib 10 in \z -> x + y + z`

How Many Parameters? 2

However, the compiler actually makes a finer grained distinction:

Two: `f2 = \y -> \z -> let x = fib 10 in x + y + z`

One: `f1 = \y -> let x = fib 10 in \z -> x + y + z`


Zero! `f0 = let x = fib 10 in \y -> \z -> x + y + z`

Arity is the number of arguments we can apply before the function does some “real work”

Arity depends on the **implementation** of the function, not its type!

Representing Functions

```
f1 = \y -> let x = fib 10 in \z -> x + y + z
```

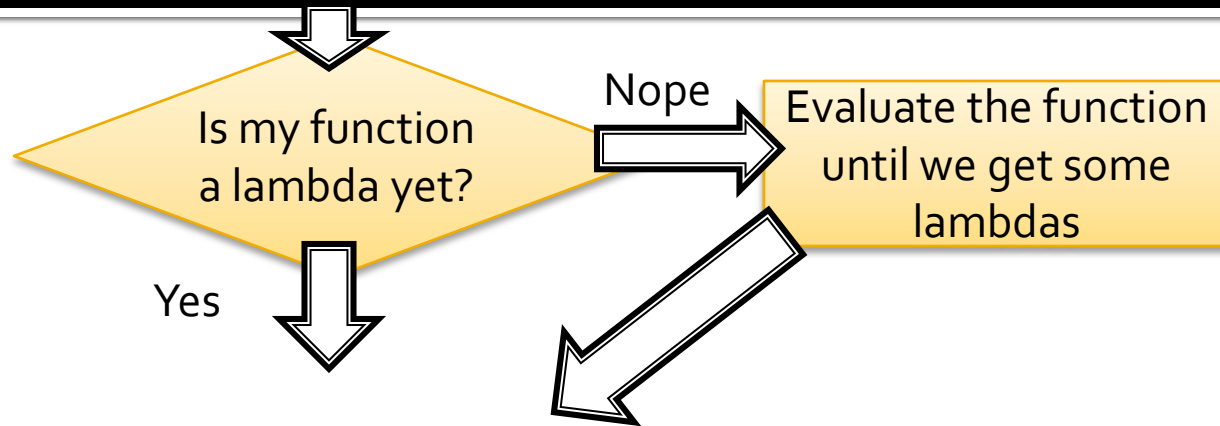


Function Application Today

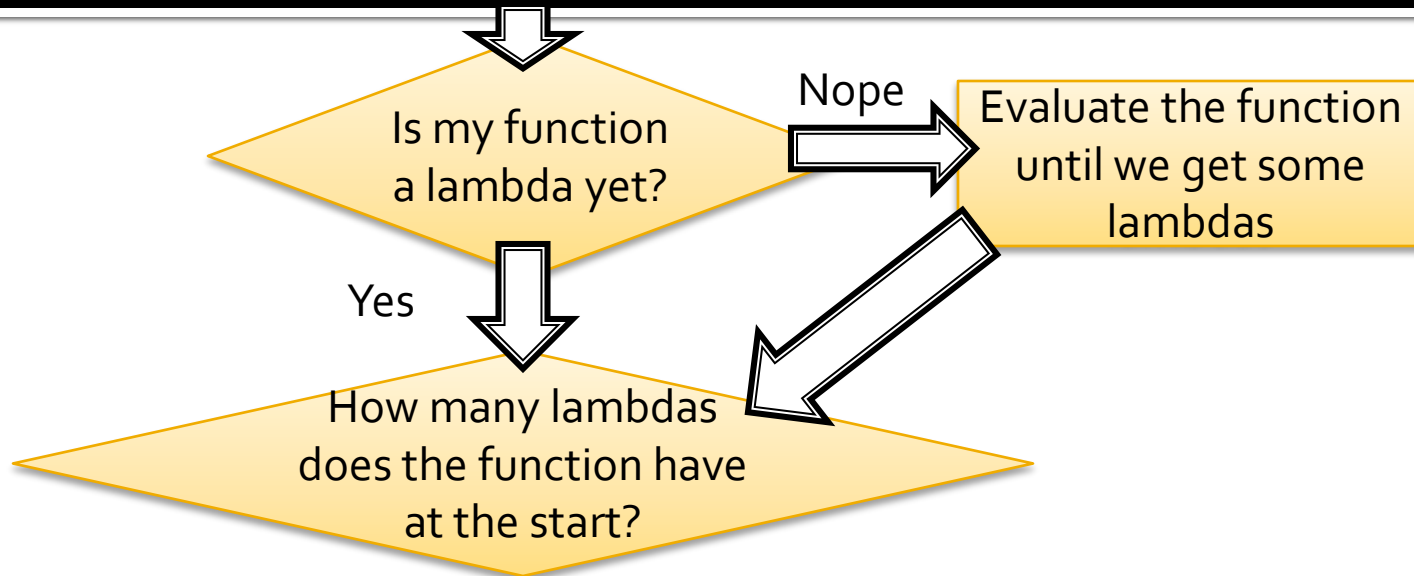


Is my function
a lambda yet?

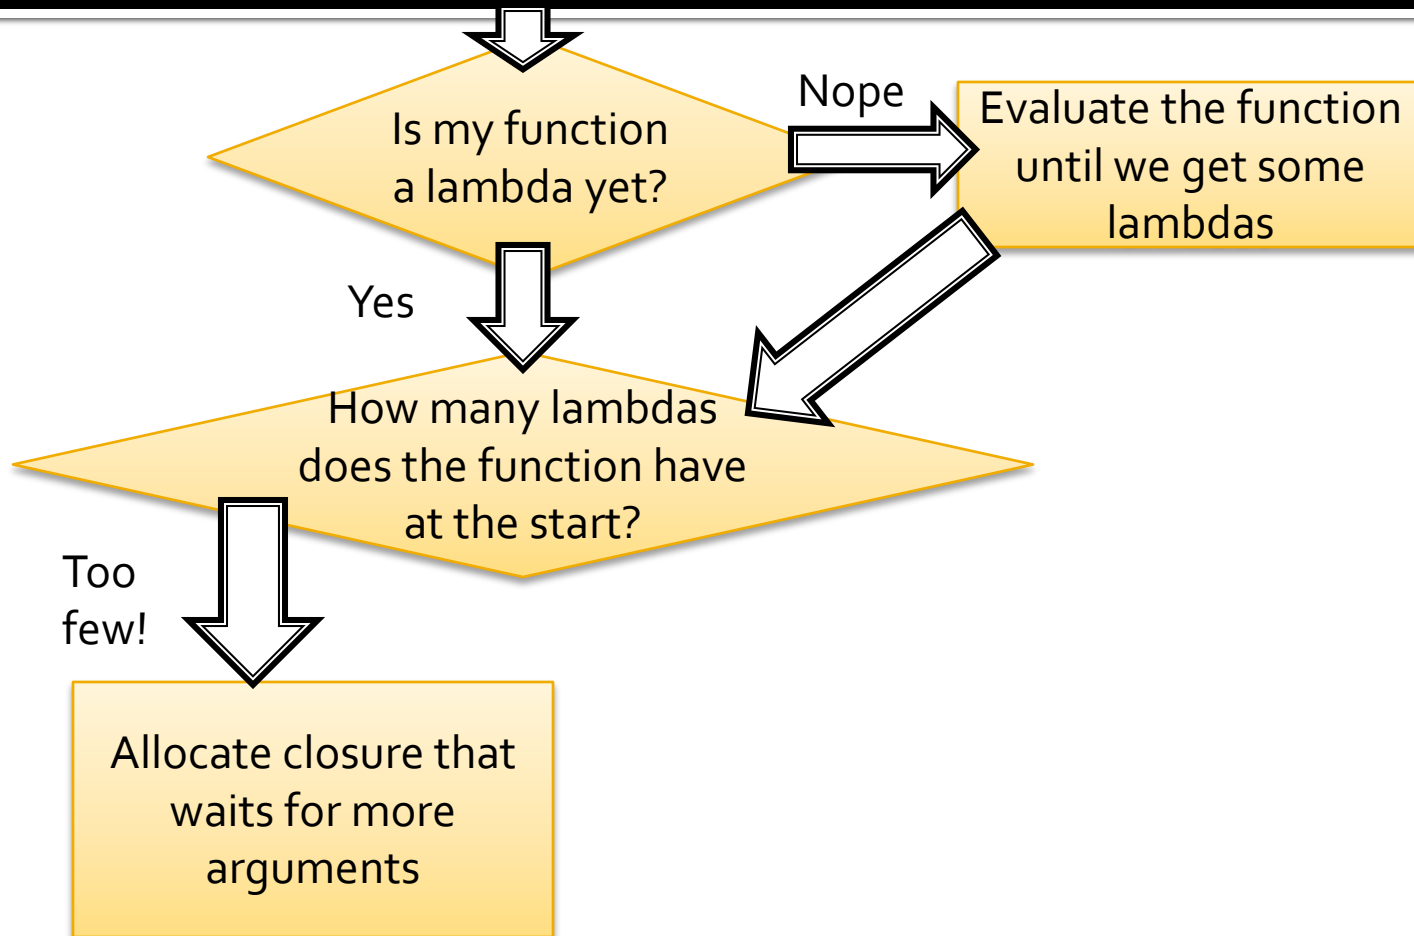
Function Application Today



Function Application Today

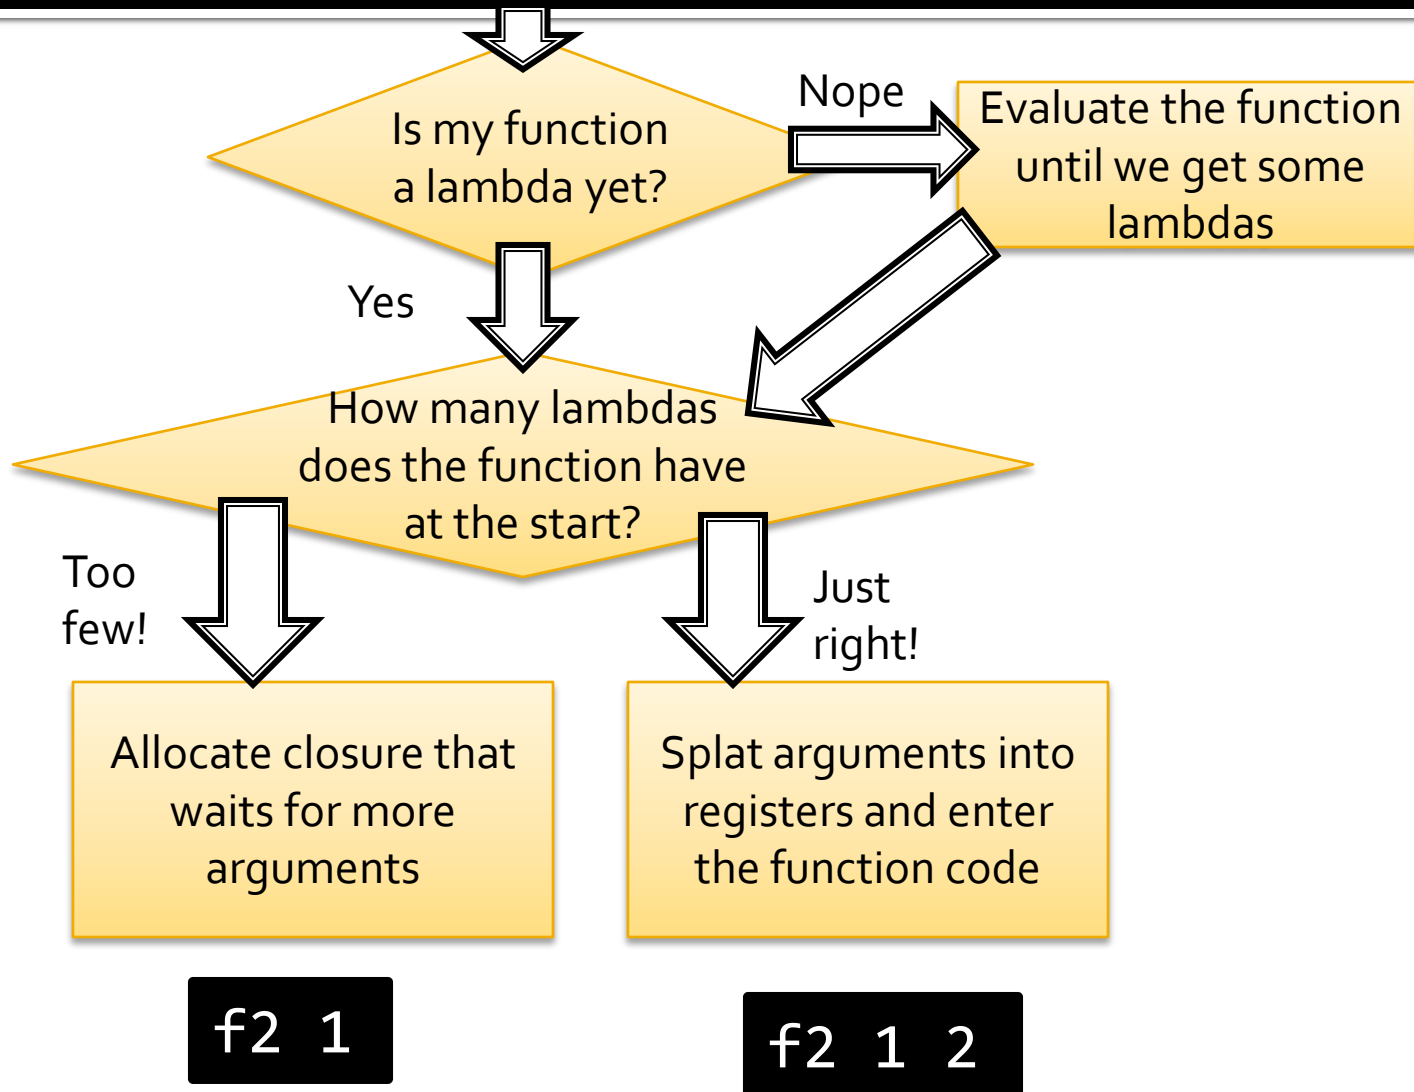


Function Application Today




f2 1

Function Application Today



Function Application Today




What Sucks

- Mandatory **dynamic** check of the function's runtime *arity* against the number of available arguments
- Mandatory **dynamic** check that the function has actually been evaluated to some lambdas
 - Things are actually worse than this: for example, there is no way to encode the fact that e.g. a `Bool` argument must have been evaluated by the caller!
 - So **unnecessary thunks are being allocated**—bad!


Multiple Results

```
silly x = (x, x + 1)  
result = silly 10
```

What we **get**:



What we **want**:



The Idea

Key idea: put juicy operational information relating to *arity*, *thunking* and *unboxing* in the type system of the compiler intermediate language

- Invariants enforced via the type system are **stable** – if you break them, your program doesn't typecheck any longer! Easy to **detect compiler bugs**
- Types are a convenient place to encode all statically known information about a value:
 - Types (and their implied invariants) easily available to the code generator
 - Seamless treatment of invariants on **higher order arguments!**

Slogan: (intermediate language) types **are** calling conventions

How We Do It: Strict Core

Binders

b	$::=$	$x:\tau$	Value binding
	$ $	$\alpha:\kappa$	Type binding

Types

τ, v, σ	$::=$	\mathbf{T}	Type constructors
	$ $	α	Type variable references
	$ $	$\bar{b} \rightarrow \bar{\tau}$	Function types
	$ $	τv	Type application

Basically: if you have n types in <brackets> to the left (right) of a function arrow then n arguments (results) are passed into (out of) the function call **in registers**

Examples

Expressing the number of **results** we expect:

```
duplicate :: <Int> -> <Int, Int>  
duplicate = \<x> -> <x, x>
```

Expressing the number of **arguments** we expect:

```
add2 :: <Int, Int> -> <Int>  
add2 = \<x, y> -> <x + y>
```

Polymorphism:

```
id :: <a :: *, a> -> <a>
```

Thunks

- Thunks are **zero-argument functions**
 - Very natural – lambdas somehow induce a **delay** in evaluation
 - Similar to classic lazy lists in ML
- However, we want to **share the result of the function** between several calls
 - Add special cases to the operational semantics that **update** thunks with what they evaluate to

 $\{\tau, \dots, \tau\}$ \triangleq $\langle \rangle \rightarrow \langle \tau, \dots, \tau \rangle$

Let's Optimize!

- I've shown you some of the calling conventions we would **like** to have for our Haskell source functions
- The straightforward translation from Haskell will **not** give you these optimized calling conventions directly
 - Instead, we need to have some optimisations that improve calling conventions through program transformation

Constructed Product Result

```
silly :: <{Int}> -> <({Int}, {Int})>  
silly = \<x>. <(x, {x <> + 1})>
```

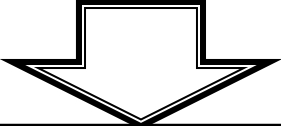
We take an argument and then **immediately** return a **product type**? We could cancel that with a use site! **Let's optimize:**

```
silly' :: <{Int}> -> <{Int}, {Int}>  
silly' = \<x>. <x, {x <> + 1}>
```

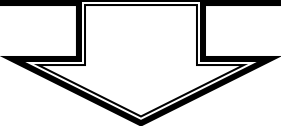
```
silly = \<x>. let <y, z> = silly' <x>  
              in <(y, z)>
```

Worker/Wrapper In Action

```
case silly <{10}> of  
  (x, y) -> <x + y>
```



```
case let <a, b> = silly' <{10}>  
  in <(a, b)> of  
  (x, y) -> <x + y>
```



```
let <a, b> = silly' <{10}>  
in <a + b>
```


Arity Definition Site Analysis 1

```
g2 :: <Int> -> <<Int> -> <Int>>  
g2 = \<x>. <\<y>. ...>
```

We take an argument and then **immediately** return a **function**? That's stupid! **Let's optimize:**

```
g2' :: <Int, Int> -> <Int>  
g2' = \<x, y>. ...
```

```
g2 = \<x>. <\<y>. <g2' <x, y>>>
```

Arity Definition Site Analysis 2

Q: Can we **always** improve the function arity like that?

```
h :: <Int> -> <<Int> -> <Int>>
```

```
h = \<x>. let <z> = fib <x>  
         in \<y>. y + z
```

```
let h' = h <10000> in h' <2> + h' <1>
```

A: **No!** If we change the type of `h` to `<Int, Int> -> <Int>` then we won't be able to **share the partial application** any longer. Result: `fib <10000>` would be **run twice!**

Arity Use Site Analysis

Q: Can we improve the arity of **higher-order arguments**?

```
h :: <<Int> -> <<Int> -> <Int>>>  
    -> <Int>  
h = \<f>. f <1> <2> + f <3> <4>
```

A: **Yes!** The function *h* only ever applies *f* to two arguments at once – and **no partial application is shared!** Let's optimize:

```
h' :: <<Int, Int> -> <Int>>  
    -> <Int>  
h' = \<f'>. f' <1, 2> + f' <3, 4>  
h = \<f>. h' <\<x,y> -> f <x> <y>>
```

Strictness

```
i :: <{Bool}> -> <Int>  
i = \<b>. case b <> of True -> <...>  
                        False -> <...>
```

We have a {thunked} argument that we **immediately evaluate**? That's stupid! **Let's optimize:**

```
i' :: <Bool> -> <Int>  
i' = \<b'>. case b' of True -> <...>  
                        False -> <...>  
  
i = \<b>. i' <b <>>
```

Deep Unboxing

```
i :: <{({Int}, {Int})}> -> <Int>
i = \<p>. case ... of
  True -> <1>
  False -> case p <> of
    (x, y) -> <x <> + y <>>
```

This is more subtle! If `i` ever evaluates `p` then it certainly evaluates both components of the pair. **Let's optimize:**

```
i' :: <{Int, Int}> -> <Int>
i' = ...
i = ...
```

Conclusions

- Identifying types with calling conventions looks like a big win, optimization-wise
 - Entirely new optimisation opportunities (arity use-site analysis, deep unboxing)
 - Existing *ad hoc* optimizations put on a sound footing
- You can have a simple, strong **operational model** and syntactic **beauty** at the **same time!**
- Surprising (to us): the best choice for the intermediate language of a compiler for a lazy language is **not** *itself* a lazy language
- Exciting direction for future work: push the ability to write strict programs into Haskell!